Claude Code — Session Opener v4

Delete this header before sending. Attach: interphase_prototype_briefing_v4.md

---

I need you to build a working iPhone webapp prototype for Interphase. Full briefing attached — read it completely before writing any code.

The key technical challenge in this build is streaming + speech working together.

The experience should feel like this: user speaks a reflection, releases the mic, a thinking indicator appears briefly, then the question begins building on screen word by word — and before it's finished writing, the voice begins speaking the first sentence. Text and speech arrive simultaneously. That's what makes it feel like a real conversation rather than a request-response cycle.

The streaming + speech architecture, precisely:

1. API call with `stream: true`
2. Process the `ReadableStream` token by token via SSE
3. Render each token to the message element as it arrives
4. On first sentence-ending punctuation (`.` `?` `!`): wait 400ms, then begin SpeechSynthesis
5. Queue subsequent sentences — play immediately, no gap between them
6. Streaming cursor (blinking gold bar) at end of text while tokens arriving
7. Cursor removed when stream completes

The 400ms pause before first speech is intentional. It makes the system feel like it's thinking before it speaks. Do not remove it.

Build in this order — do not skip ahead:

1. Shell, fonts, colours, PWA meta
2. Conversation layout
3. Streaming API call — tokens rendering on screen
4. SpeechSynthesis with sentence detection — speak as stream arrives
5. Test this on a real iPhone before continuing — streaming + speech must work on device
6. SpeechRecognition hold-to-speak
7. Text input fallback
8. Simulated conversation (typewriter + speech)
9. API key modal
10. Onboarding (2 cards)

11. Chaperone slide-up panel
12. Micro-interactions

Stop at step 5 and confirm it works before building the rest. Streaming to screen is straightforward. Streaming + real-time speech triggering is the hard part. Get that right first.

Parameters that must not change:

- `max_tokens: 120` — forces short questions, do not increase
- Voice rate: `0.88`, pitch: `0.95` — warm and unhurried, not default
- Voice priority: Samantha → Daniel → en-GB → en-US → default
- 400ms delay before first sentence only — subsequent sentences queue immediately
- System prompt rule 15: AI must front-load the question, no preamble — because the first words spoken will be whatever arrives first in the stream

One question before you start: do you have access to a real iPhone for testing, or will you be testing in a desktop browser? The answer changes how we approach the SpeechRecognition implementation.